

Machine Problem 3 Report

Presented to Prof. Ahmad Afsahi
ELEC 374: Digital Systems Engineering
Faculty of Applied Science
Queen's University

Nathan Goodman: 20228249

Date: April 5th, 2023

CUDA Code

For the three different requested outputs, three different CUDA scripts were created:

1. Find the time it takes to transfer two matrices worth of data to and from device:

Note: Block width is equal to 16 and gpu execution time includes transfer times.

```
1  #include "cuda_runtime.h"
2  #include <string.h>
3  #include <stdio.h>
4  #include <stdlib.h>
5  #include <time.h>
6  #include <math.h>
7
8  #define debug 0 //If debug = 1, certain print statements will be enabled
9
10 /*
11 1. Code __global__ (GPU) function for matrixMultiplication
12 */
13 __global__ void matrixMultiplication_kernel(float*d_a, float*d_b, float*d_c, int sizeofMatricies) {
14     int Row = blockIdx.y*blockDim.y + threadIdx.y;
15     int Col = blockIdx.x*blockDim.x + threadIdx.x;
16     if (Row < sizeofMatricies && Col < sizeofMatricies) {
17         float Pvalue = 0;
18         for (int k = 0; k < sizeofMatricies; k++) {
19             Pvalue += d_a[Row*sizeofMatricies + k] * d_b[k*sizeofMatricies + Col];
20         }
21         d_c[Row*sizeofMatricies + Col] = Pvalue;
22     }
23 }
24
25 /*
26 1b. Create regular C function (to run on CPU) for matrixMultiplication
27 */
28 void matrixMultiplication(float *a, float *b, float *c, int sizeofMatricies) {
29     for (int i = 0; i < sizeofMatricies; i++) { //i =row
30         for (int j = 0; j < sizeofMatricies; j++) { //j =col
31             //For each value of matrix c (c[i][j]):
32             float sum = 0;
33             for (int k = 0; k < sizeofMatricies; k++) {
34                 sum += *(a + i*sizeofMatricies + k) * *(b + k*sizeofMatricies + j);
35                 //sum = sum + a[Row][i] * b[i][Col]
36             }
37             *(c + i*sizeofMatricies + j) = sum;
38         }
39     }
40 }
41
42 //Helpful Functions from MP2:
43 void fillMatrix(float *a, int size) {
44     for (int i = 0; i < size; i++) {
45         for (int j = 0; j < size; j++) {
46             *(a + i*size + j) = rand() % 100; //Every element will be a random number in range of 0 to 100
47         }
48     }
49 }
50
51 int correct_output(float *a, float *b, int sizeofMatricies) {
52     //Checks if two matrices are equal returns 1 if yes 0 if no.
53     for (int i = 0; i < sizeofMatricies; i++) {
54         for (int j = 0; j < sizeofMatricies; j++) {
55             if (*(a + i*sizeofMatricies + j) != *(b + i*sizeofMatricies + j)) {
56                 //If a[i][j] != b[i][j]:
57                 return 0;
58             }
59         }
60     }
61     return 1;
62 }
63 void printMatrix(float *a, int size) {
64     for (int i = 0; i < size; i++) {
65         for (int j = 0; j < size; j++) {
66             printf("%f ", *(a + i*size + j));
67         }
68         printf("\n");
69     }
70 }
```

```

72 int main(int argc, char *argv[]) {
73     cudaDeviceProp deviceProps;
74
75     //Get Device Name (They did this in tutorial so I'll do it here)
76     cudaGetDeviceProperties(&deviceProps, 0);
77     printf("CUDA device [%s]\n", deviceProps.name);
78     if (debug == 1) {
79         printf("\tNumber of Multprocessors: %d\n", deviceProps.multiProcessorCount);
80         printf("\tMax Threads Per Block: %d\n", deviceProps.maxThreadsPerBlock);
81         printf("\tMax Dimension of a Block: %d\n", deviceProps.maxThreadsDim);
82         printf("\tMax Dimension of a Grid: %d\n", deviceProps.maxGridSize);
83     }
84
85     //Set Size of Matrices:
86     const int dimOfMatrices = 2000; //Value determines size of matrices Ex. dimOfMatrices = 5 will result in 5x5 matrices
87     printf("Dimensions of Matrices: %dx%d\n", dimOfMatrices, dimOfMatrices);
88
89     //Calculate amount of memory they take:
90     int nbytes = dimOfMatrices*dimOfMatrices*sizeof(float);
91
92     //Allocate host memory for matrices:
93     float *a = 0;
94     float *b = 0;
95     float *c = 0;
96     float *cCopy = 0; //Pointer to a copy of second output matrix, to be used to observe GPU to host transfer time
97     cudaMallocHost((void*)&a, nbytes); //Allocates host memory for matrix A, and points pointer a to first value.
98     cudaMallocHost((void*)&b, nbytes);
99     cudaMallocHost((void*)&c, nbytes);
100     cudaMallocHost((void*)&cCopy, nbytes);
101
102     if (debug == 1) {
103         printf("Memory Locations of matrices:\n");
104         printf("\ta = %x\n", &a);
105         printf("\tb = %x\n", &b);
106         printf("\tc = %x\n", &c);
107     }
108
109     //Generate input matrices into memory
110
111     srand(time(NULL));
112     fillMatrix(a, dimOfMatrices);
113     fillMatrix(b, dimOfMatrices);
114
115     if (debug == 1) {
116         printf("Matrix A:\n");
117         printMatrix(a, dimOfMatrices);
118         printf("Matrix B:\n");
119         printMatrix(b, dimOfMatrices);
120     }
121
122     //1e. Allocate device memory for matrices:
123     float *d_a = 0;
124     float *d_b = 0;
125     float *d_c = 0;
126     cudaMalloc((void*)&d_a, nbytes); //Allocates memory for matrix A, and points pointer a to first value.
127     cudaMemcpy(d_a, 255, nbytes); //Sets all allocated bytes to 255 (they did this in tutorial so i did it here)
128     cudaMalloc((void*)&d_b, nbytes);
129     cudaMemcpy(d_b, 255, nbytes);
130     cudaMalloc((void*)&d_c, nbytes);
131     cudaMemcpy(d_c, 255, nbytes);
132
133     //Set kernel launch configuration
134     int blkWidth = 16;
135     int threadsPerBlock = blkWidth*blkWidth;
136     int threadsNeeded = ceil(dimOfMatrices*dimOfMatrices); //Because in this configuration one thread only produces one value
137     int blocksNeeded = ceil(threadsNeeded / threadsPerBlock);
138     if (blocksNeeded < 1) blocksNeeded++;
139     int gridWidth = ceil(sqrt(blocksNeeded)); //Note grid will be square

```

```

141     if (debug == 1) {
142         printf("Block Width: %d\n", blkWidth);
143         printf("Grid Width: %d\n", gridWidth);
144     }
145
146     dim3 dimBlocks = dim3(blkWidth, blkWidth);
147     dim3 dimGrid = dim3(gridWidth, gridWidth);
148
149     //create cuda event handles
150     cudaEvent_t start, stopToDevice, startToHost, stop;
151     cudaEventCreate(&start);
152     cudaEventCreate(&stopToDevice);
153     cudaEventCreate(&startToHost);
154     cudaEventCreate(&stop);
155
156     cudaDeviceSynchronize();
157     float gpu_time = 0.0f;
158     float toGPU_time = 0.0f;
159     float fromGPU_time = 0.0f;
160
161     //asynchronously issue work to the GPU (all stream 0)
162     cudaEventRecord(start, 0);
163
164     //Copy inputs to device
165     cudaMemcpyAsync(d_a, a, nbytes, cudaMemcpyHostToDevice, 0);
166     cudaMemcpyAsync(d_b, b, nbytes, cudaMemcpyHostToDevice, 0);
167
168     //Save time it took to write to device
169     cudaEventRecord(stopToDevice, 0);
170
171     //Call Kernel
172     matrixMultiplication_kernel << <dimGrid, dimBlocks, 0, 0 >> >(d_a, d_b, d_c, dimOfMatrices);
173     //Note: only section after >>> is the actual function parameters
174
175     //Save time before saving to host
176     cudaEventRecord(startToHost, 0);
177
178     //Copy outputs from device
179     cudaMemcpyAsync(c, d_c, nbytes, cudaMemcpyDeviceToHost, 0);
180     cudaMemcpyAsync(cCopy, d_c, nbytes, cudaMemcpyDeviceToHost, 0); //A second matrix is recorded as question asked for transfer to two matrices
181
182     cudaEventRecord(stop, 0);
183     cudaEventSynchronize(stopToDevice);
184     cudaEventSynchronize(startToHost);
185     cudaEventSynchronize(stop);
186     cudaEventElapsedTime(&gpu_time, start, stop); //Note: gpu_time includes toGPU_time and fromGPU_time
187     cudaEventElapsedTime(&toGPU_time, start, stopToDevice);
188     cudaEventElapsedTime(&fromGPU_time, startToHost, stop);
189
190     //print the GPU times
191     printf("time spent executing by the GPU: %.2f\n", gpu_time);
192     printf("time spent transferring input data to the GPU: %.2f\n", toGPU_time);
193     printf("time spent transferring output data to host: %.2f\n", fromGPU_time);
194
195     //Calculate matrixMultiplication using CPU:
196     float *d = 0;
197     cudaMallocHost((void**)&d, nbytes);
198     matrixMultiplication(a, b, d, dimOfMatrices); //d is output matrix
199
200     //Check Output
201     bool bFinalResults = (bool)correct_output(c, d, dimOfMatrices); //check if c(from GPU) = d(from CPU)
202     if (bFinalResults == true) {
203         printf("Test PASSED\n");
204     }
205
206     if (debug == 1) {
207         printf("Matrix C:\n");
208         printMatrix(c, dimOfMatrices);
209         printf("Correct Matrix:\n");
210         printMatrix(d, dimOfMatrices);
211     }
212
213     //release resources
214     cudaEventDestroy(start);
215     cudaEventDestroy(stop);
216
217     cudaFreeHost(a);
218     cudaFreeHost(b);
219     cudaFreeHost(c);
220     cudaFreeHost(cCopy);
221     cudaFreeHost(d);
222     cudaFree(d_a);
223     cudaFree(d_b);
224     cudaFree(d_c);
225     cudaDeviceReset();
226
227     return 0;
228
229 }

```

2. Find time it takes to execute matrix on GPU with one block and one thread vs CPU

Note: block width is equal to 16. gpu execution time does NOT include transfer times.

```

1  #include "cuda_runtime.h"
2  #include <string.h>
3  #include <stdio.h>
4  #include <stdlib.h>
5  #include <time.h>
6  #include <math.h>
7
8  #define debug 0 //If debug = 1, certain print statements will be enabled
9
10 /*
11 2. Code __global__(GPU) function for matrixMultiplication but one thread calculates whole output matrix
12 */
13 __global__ void matrixMultiplication_kernel(float*d_a, float*d_b, float*d_c, int sizeOfMatrices) {
14     for (int i = 0; i < sizeOfMatrices; i++) { //i =row
15         for (int j = 0; j < sizeOfMatrices; j++) { //j =col
16             float Pvalue = 0;
17             for (int k = 0; k < sizeOfMatrices; k++) {
18                 Pvalue += d_a[i*sizeOfMatrices + k] * d_b[k*sizeOfMatrices + j];
19             }
20             d_c[i*sizeOfMatrices + j] = Pvalue;
21         }
22     }
23 }
24
25 //Regular C function (to run on CPU) for matrixMultiplication:
26 void matrixMultiplication(float *a, float *b, float *c, int sizeOfMatrices) {
27     for (int i = 0; i < sizeOfMatrices; i++) { //i =row
28         for (int j = 0; j < sizeOfMatrices; j++) { //j =col
29             //For each value of matrix c (c[i][j]):
30             float sum = 0;
31             for (int k = 0; k < sizeOfMatrices; k++) {
32                 sum += *(a + i*sizeOfMatrices + k) * *(b + k*sizeOfMatrices + j);
33                 //sum = sum + (a[Row][i] * b[i][Col])
34             }
35             *(c + i*sizeOfMatrices + j) = sum;
36         }
37     }
38 }
39
40 //Helpful Functions from MP2:
41
131 //Set kernel launch configuration
132 int blkWidth = 1;
133 int gridWidth = 1;
134 /* Below is not needed for this part as it asks for one block and one thread
135 int threadsPerBlock = blkWidth*blkWidth;
136 int threadsNeeded = ceil(dimOfMatrices*dimOfMatrices); //Because in this configuration one thread only produces one value
137 int blocksNeeded = ceil(threadsNeeded / threadsPerBlock);
138 if (blocksNeeded < 1) blocksNeeded++;
139 int gridWidth = ceil(sqrt(blocksNeeded)); //Note grid will be square
140 */
141
142 //print the GPU times
143 printf("time spent executing by the GPU: %.2f\n", (gpu_time - toGPU_time - fromGPU_time)); //Time spent executing by the GPU not including transfers
144 printf("time spent transferring input data to the GPU: %.2f\n", toGPU_time);
145 printf("time spent transferring output data to host: %.2f\n", fromGPU_time);

```

Note: All parts not shown were identical to the parts from the first configuration. Full script can be viewed from attached .zip

3. Increase block_width as matrix size increases.

Note: Gpu execution time does NOT include transfer times.

```

138 //Set kernel launch configuration
139 int blkWidth = 25;
140 int threadsPerBlock = blkWidth*blkWidth;
141 int threadsNeeded = ceil(dimOfMatrices*dimOfMatrices); //Because in this configuration one thread only produces one value
142 int blocksNeeded = ceil(threadsNeeded / threadsPerBlock);
143 if (blocksNeeded < 1) blocksNeeded++;
144 int gridWidth = ceil(sqrt(blocksNeeded)); //Note: grid will be square
145
146 printf("Block Width: %d\n", blkWidth);
147 printf("Grid Width: %d\n", gridWidth);
148

```

```

191 //print the GPU times
192 printf("time spent executing by the GPU: %.2f\n", (gpu_time - toGPU_time - fromGPU_time)); //Time spent executing by the GPU not including transfers
193 //printf("time spent transferring input data to the GPU: %.2f\n", toGPU_time);
194 //printf("time spent transferring output data to host: %.2f\n", fromGPU_time);

```

Note: All parts not shown were identical to the parts from the first configuration. Full scripts can be viewed from attached .zip

Output

For each question, outputs were recorded for the following matrix sizes: 125x125, 250x250, 500x500, 1000x1000, 2000x2000; with each configuration from above.

1. Find the time it takes to transfer two matrices worth of data to and from device:

```

CUDA device [Tesla C2075]
Dimensions of Matrices: 125x125
time spent executing by the GPU: 0.28
time spent transferring input data to the GPU: 0.08
time spent transferring output data to host: 0.03
Test PASSED
Press any key to continue . . . _

```

```

CUDA device [Tesla C2075]
Dimensions of Matrices: 250x250
time spent executing by the GPU: 1.20
time spent transferring input data to the GPU: 0.13
time spent transferring output data to host: 0.08
Test PASSED
Press any key to continue . . . _

```

```

CUDA device [Tesla C2075]
Dimensions of Matrices: 500x500
time spent executing by the GPU: 7.75
time spent transferring input data to the GPU: 0.38
time spent transferring output data to host: 0.32
Test PASSED
Press any key to continue . . . _

```

```

CUDA device [Tesla C2075]
Dimensions of Matrices: 1000x1000
time spent executing by the GPU: 58.40
time spent transferring input data to the GPU: 1.35
time spent transferring output data to host: 1.24
Test PASSED
Press any key to continue . . . _

```

```

CUDA device [Tesla C2075]
Dimensions of Matrices: 2000x2000
time spent executing by the GPU: 444.63
time spent transferring input data to the GPU: 5.27
time spent transferring output data to host: 4.92
Test PASSED
Press any key to continue . . . _

```

2. Find time it takes to execute matrix on GPU with one block and one thread vs CPU

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 125x125
time spent executing by the GPU: 811.50
time spent executing by CPU: 10.00
Test PASSED
Press any key to continue . . . _
```

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 250x250
time spent executing by the GPU: 9092.94
time spent executing by CPU: 106.00
Test PASSED
Press any key to continue . . . _
```

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 500x500
time spent executing by the GPU: 73857.30
time spent executing by CPU: 1001.00
Test PASSED
Press any key to continue . . . _
```

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 1000x1000
time spent executing by the GPU: 593185.50
time spent executing by CPU: 12923.00
Test PASSED
Press any key to continue . . . _
```

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 2000x2000
time spent executing by the GPU: 4745166.00
time spent executing by CPU: 119458.00
Test PASSED
Press any key to continue . . . _
```

3. Increase block_width as matrix size increases.

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 125x125
Block Width: 2
Grid Width: 63
time spent executing by the GPU: 2.86
Test PASSED
Press any key to continue . . . _
```

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 250x250
Block Width: 4
Grid Width: 63
time spent executing by the GPU: 6.13
Test PASSED
Press any key to continue . . . _
```

```
CUDA device [Tesla C2075]
Dimensions of Matrices: 500x500
Block Width: 10
Grid Width: 50
time spent executing by the GPU: 9.00
Test PASSED
Press any key to continue . . . _
```

```

CUDA device [Tesla C2075]
Dimensions of Matrices: 1000x1000
Block Width: 20
Grid Width: 50
time spent executing by the GPU: 62.11
Test PASSED
Press any key to continue . . . _

```

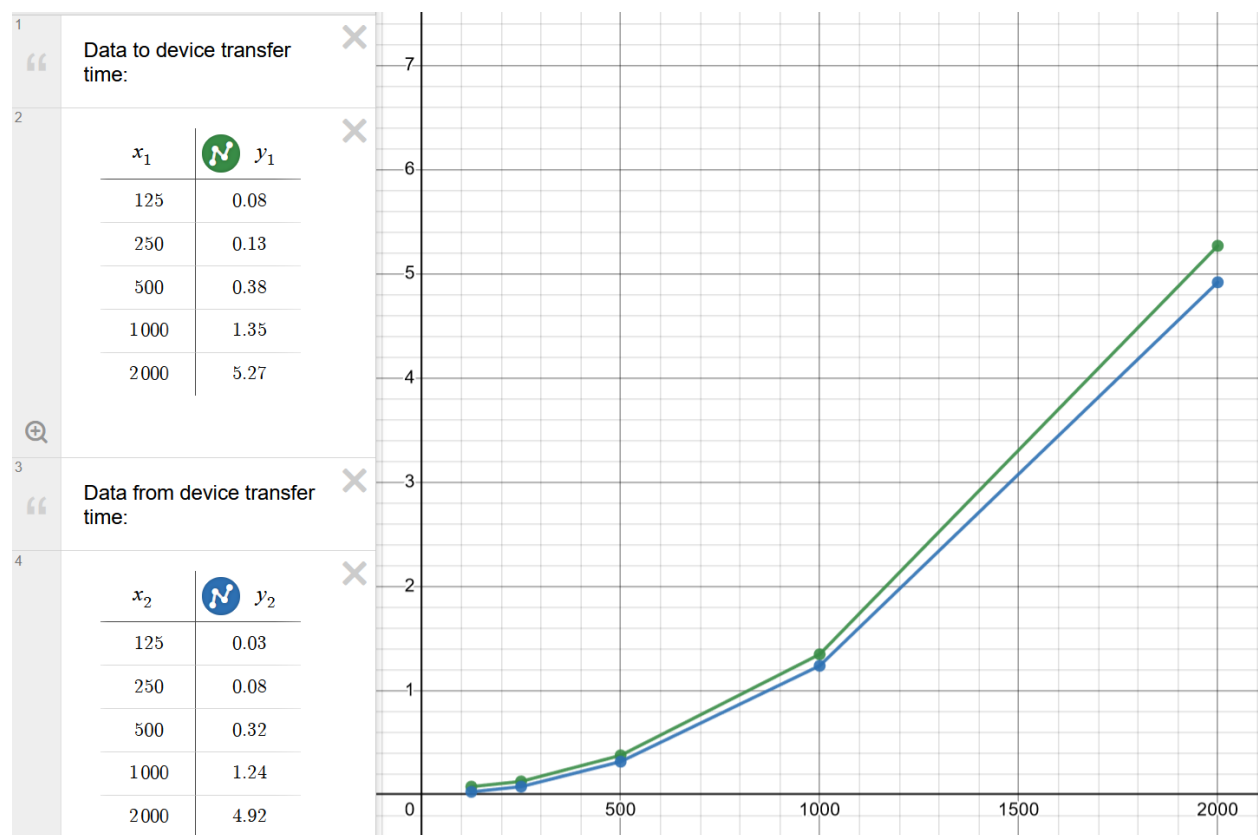
```

CUDA device [Tesla C2075]
Dimensions of Matrices: 2000x2000
Block Width: 25
Grid Width: 80
time spent executing by the GPU: 516.63
Test PASSED
Press any key to continue . . . _

```

Analysis

1. For the first configuration, the following graph showing transfer times to and from device versus matrix size, were generated from the outputs above. The x-axis is block width and the y-axis is transfer time (in milliseconds):

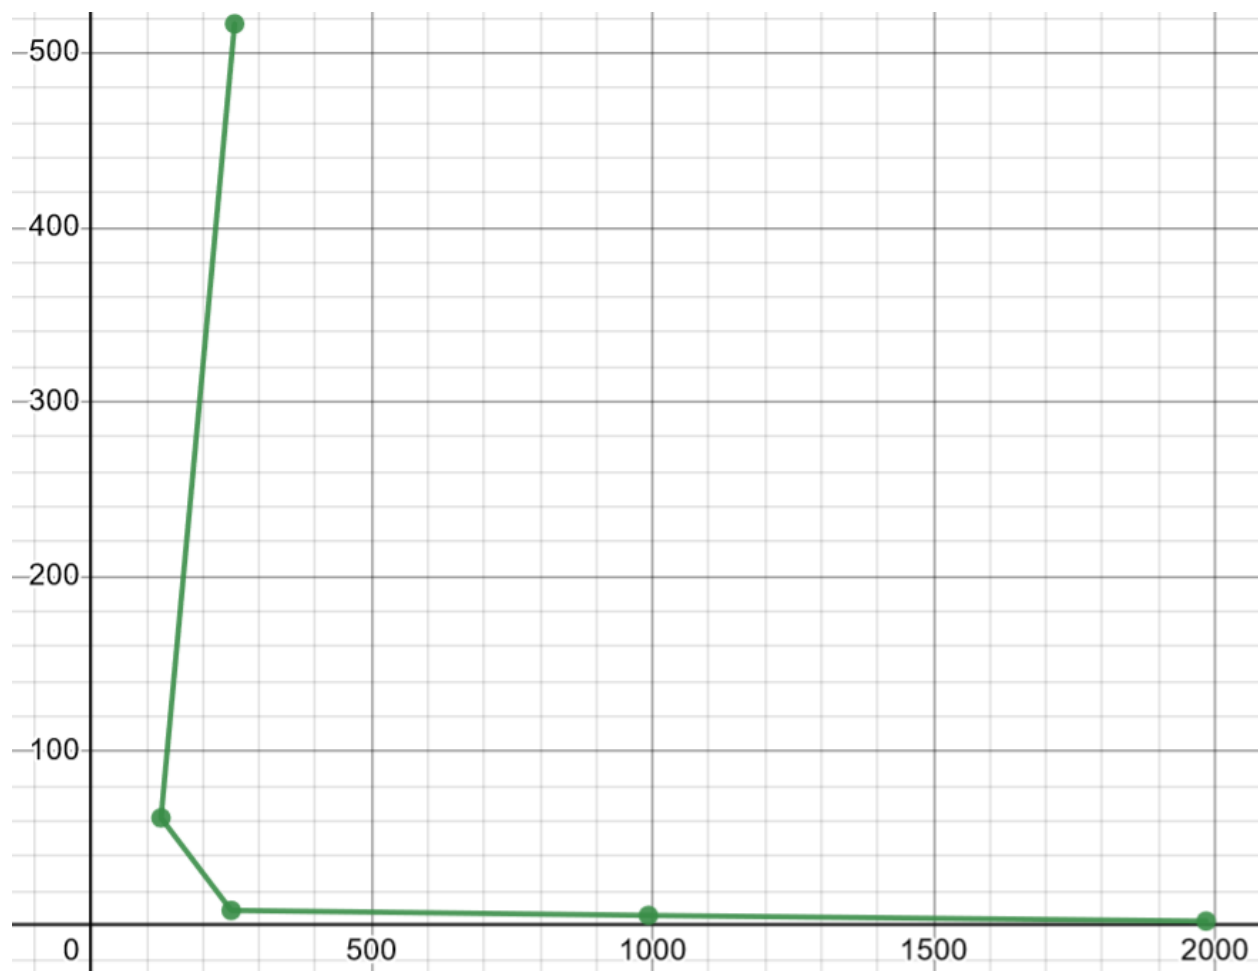


There is for the most part no difference in transfer speeds to and from the gpu device, with only a slight—seemingly proportional difference being seen (difference < 0.4 ms for 2000x2000). This small difference is likely due to a difference in memory speeds between the device and host.

2. Observing the outputs from the second configuration (above in Outputs), it becomes obvious that setting the grid width and block width to one had a detrimental impact on the GPU's performance. When comparing this performance to the performance of the matrix multiplication on the CPU, it's clear that it is not always beneficial to offload your matrix multiplication to the GPU; it is only optimal to do so if the kernel configuration is optimized correctly.

3. The outputs for the third configuration were generated from the following matrixWidth and blockWidth pairs: (125, 2), (250, 4), (500, 10), (1000, 20), (2000, 25).

The following graph of kernel execution time vs (numBlocks/blockWidth) were generated from the outputs above (Note: numBlocks = gridWidth²). The x-axis is numBlocks/blockWidth and the y-axis is execution time (in milliseconds):



This clearly shows that when there are more blocks with less threads per block, the execution time will be faster than when there is less blocks with more threads per block.

a) Using our knowledge of the matrix multiplication algorithm, we know that to calculate an element in output matrix C, we must access every element from input matrix A sharing a row with the output element, and we must access every element from input matrix B sharing a column with the output element. This means for every element in matrix C we are accessing $2 \times \text{matrixWidth}$ number of input elements:

number of total accesses = amount of output matrix elements * 2 * matrixWidth

Since C is a square matrix 'matrixWidth' wide, the number of output elements will be: matrixWidth²

If we are trying to find the amount of times each input element is accessed:

$$\text{average number of times an input element is accessed} = \frac{\text{number of total accesses}}{\text{number of input elements}}$$

The number of input elements is equal to all the elements in A plus all the elements from B. Since both of these square matrices are a 'matrixWidth' wide:

$$\text{Number of input elements} = 2(\text{matrixWidth}^2)$$

Combining all of these into one equation:

$$\begin{aligned} \text{avgInputElementAccesses} &= \frac{\text{totalAccesses}}{\text{numberOfInputElements}} \\ \text{avgInputElementAccesses} &= \frac{\text{matrixWidth}^2 * (2 * \text{matrixWidth})}{2 * \text{matrixWidth}^2} \\ \text{avgInputElementAccesses} &= \text{matrixWidth} \end{aligned}$$

Therefore, for this configuration where the output matrixWidth is equal to the input matrices matrixWidth, the number of times a given input matrix element will be accessed is equal to the matrix width. This makes sense as a given element in A or B needs to be accessed once for each element in C, in the same row/column (=matrixWidth).

b) CGMA ratio is the number of operations a thread performs over the number of memory fetches.

Looking at our matrix multiplication kernel for the third configuration (identical to the one for the first configuration), for each loop iteration within the kernel:

```
for (int k = 0; k < sizeofMatrices; k++) {  
    Pvalue += d_a[Row*sizeofMatrices + k] * d_b[k*sizeofMatrices + Col];  
}
```

There's two global memory accesses d_a[...] and d_b[...] for one floating-point multiplication and one floating-point addition operation. Giving a CGMA ratio of floating-point operations to bytes accessed from global memory of 0.25 FLOP/B.